

Moore-Neighbor Tracing

Idea

The idea behind Moore-Neighbor tracing is simple; but before we explain it, we need to define an important concept: the **Moore neighborhood** of a pixel.

Moore Neighborhood

The Moore neighborhood of a pixel, **P**, is the set of 8 pixels which share a vertex or edge with that pixel. These pixels are namely pixels **P1, P2, P3, P4, P5, P6, P7 and P8** shown in **Figure 1** below. The Moore neighborhood (also known as the **8-neighbors** or **indirect neighbors**) is an important concept that frequently arises in the literature.




Figure 1

Now we are ready to introduce the idea behind Moore-Neighbor tracing...

Given a digital pattern i.e. a group of black pixels, on a background of white pixels i.e. a grid; locate a black pixel and declare it as your "**start**" pixel. (Locating a "**start**" pixel can be done in a number of ways; we'll start at the bottom left corner of the grid, scan each column of pixels from the bottom going upwards - starting from the leftmost column and proceeding to the right- until we encounter a black pixel. We'll declare that pixel as our "**start**" pixel.)

Now, imagine that you are a bug (ladybird) standing on the **start** pixel as in **Figure 2** below. Without loss of generality, we will extract the contour by going around the pattern in a clockwise direction. (It doesn't matter which direction you choose as long as you stick with your choice throughout the algorithm).

The general idea is: every time you hit a black pixel, **P**, backtrack i.e. go back to the white pixel you were previously standing on, then, go **around** pixel **P** in a clockwise direction, visiting each pixel in its Moore neighborhood, until you hit a black pixel. The algorithm terminates when the start pixel is visited for a second time.

The black pixels you walked over will be the contour of the pattern.




Figure 2

Algorithm

The following is a formal description of the Moore-Neighbor tracing algorithm:

Input: A square tessellation, T , containing a connected component P of black cells.

Output: A sequence $B(b_1, b_2, \dots, b_k)$ of boundary pixels i.e. the contour.

Define $M(a)$ to be the Moore neighborhood of pixel a .

Let p denote the current boundary pixel.

Let c denote the current pixel under consideration i.e. c is in $M(p)$.

Begin

- Set B to be empty.
- From bottom to top and left to right scan the cells of T until a black pixel, s , of P is found.
- Insert s in B .
- Set the current boundary point p to s i.e. $p=s$
- Backtrack i.e. move to the pixel from which s was entered.
- Set c to be the next clockwise pixel in $M(p)$.
- While c not equal to s do

If c is black

- insert c in B
- set $p=c$
- backtrack (move the current pixel c to the pixel from which p was entered)

else

- advance the current pixel c to the next clockwise pixel in $M(p)$

end While

End

Demonstration

The following is an animated demonstration of how Moore-Neighbor tracing proceeds to trace the contour of a given pattern.

(We have decided to trace the contour in a clockwise direction).

Moore's Algorithm Demonstration



Analysis

The main weakness of Moore-Neighbor tracing lies in the choice of the stopping criterion, in other words, when does the algorithm terminate?

In the original description of the algorithm used in Moore-Neighbor tracing, the stopping criterion is visiting the **start** pixel for a second time. Like in the case of the [Square Tracing algorithm](#), it turns out that Moore-Neighbor tracing will fail to contour trace a large family of patterns if it were to depend on that criterion. What follows is an animated demonstration explaining how Moore-Neighbor tracing fails to extract the contour of a pattern due to the bad choice of the stopping criterion:

Demonstration:

**A reason to change
the
stopping criterion**

As you can see, improving the stopping criterion would be a good start to improving the overall performance of Moore-Neighbor tracing. There are 2 effective alternatives to the existing stopping criterion:

- a) Stop after visiting the **start** pixel n times, where n is at least 2, OR
- b) Stop after entering the **start** pixel a second time **in the same manner you entered it initially**. This criterion was proposed by [Jacob Eliosoff](#) and we will therefore call it **Jacob's stopping criterion**.

Using Jacob's stopping criterion will greatly improve the performance of Moore-Neighbor tracing making it the best algorithm for extracting the contour of any pattern no matter what its [connectivity](#).

The reason for this is largely due to the fact that the algorithm checks the whole [Moore neighborhood](#) of a boundary pixel in order to find the next boundary pixel. Unlike the [Square Tracing algorithm](#), which makes either left or right turns and misses "diagonal" pixels; Moore-Neighbor tracing will always be able to extract the outer boundary of any connected component. The reason for that is: for any [8-connected](#) (or simply **connected**) pattern, the **next** boundary pixel lies within the Moore neighborhood of the current boundary pixel. Since Moore-Neighbor tracing proceeds to check every pixel in the Moore neighborhood of the current boundary pixel, it is bound to detect the next boundary pixel.

When Moore-Neighbor tracing visits the start pixel for a second time in the same way it did the first time around, this means that it has traced the **complete outer contour** of the pattern and if not terminated, it will trace the same contour again. This result has yet to be proved...

All comments and questions are welcomed...
abeer.ghuneim@mail.mcgill.ca